

ML Breadth Knowledge

lzdh

November 2025

About this document

The content of this file come from various sources, mainly generative AI tools and Chip Huyen's (<https://huyenchip.com>) Machine Learning Interviews book (with deep appreciation).

It was created during the author's own interview process and is intended solely for learning and personal study. It may contain mistakes or outdated information, and accuracy is not guaranteed. Please independently verify any technical concepts before relying on them.

This material does not fully reflect the depth and breath of the real ML interviews. Some topics have been expanded with more in-depth content based on the author's personal interests, while others have been intentionally simplified for quicker review.

The author is actively updating the document and fixing the errors. If you spot any issues or have suggestions, please feel free to submit a pull request on the author's GitHub repository.

Finally, enjoy reading and enjoy prep! Good luck!

About the author

Personal Website: https://lzdh.github.io/my_website/

More interview prep notes can be found on my Rednote (xiaohongshu). Please follow!

Rednote ID: 612410210 Jessie

Github: @lzdh

LinkedIn: liu-zdh

Contents

I	ML, Stats, ...	7
1	Basics	8
1.1	Overfitting vs Underfitting	8
1.1.1	Variance–Bias Tradeoff	8
1.1.2	Solution to overfitting	8
1.2	Precision and Recall	9
1.3	Ensemble Methods	9
1.3.1	Bagging:	9
1.3.2	Boosting:	9
1.3.3	Stacking Ensemble	9
1.3.4	Cascading Ensemble	10
1.4	Decision Tree	10
1.5	Pruning of Decision Trees	10
1.6	Entropy, Gini Impurity, and Cross Entropy	10
1.7	Random Forest	11
1.8	XGBoost vs LightGBM	11
1.8.1	XGBoost	11
1.8.2	LightGBM	11
1.8.3	Comparison of Models	11
1.9	Logistic Regression	12
1.10	Logistic Regression Assumptions	12
1.11	Linear Regression Assumptions	12
1.12	K-means Assumptions	12
1.13	Feature Selection	13
1.14	Generative vs Discriminative Models	13
1.15	KL Divergence	13
1.16	L1 vs L2 Regularization	13
1.17	L2-norm vs Squared L2-norm	14
1.18	Likelihood & Posterior	14
1.19	Concept Drift (sys design)	14
1.20	Semi-Supervised Learning	14

1.21 Bayes' Rule	14
1.22 Naïve Bayes	15
1.22.1 Naive Bayes Assumptions	15
1.23 AUROC (Receiver Operating Characteristic)	15
1.24 R^2 vs Adjusted R^2	15
1.25 Correlation and Covariance	16
1.26 Maximum Likelihood Estimation (MLE)	16
1.27 Class Imbalance Loss	16
1.27.1 Re-weighted Cross Entropy:	16
1.27.2 Balanced Softmax Cross Entropy (multi-class)	17
1.27.3 Focal Loss (binary, multi-class, multi-label)	17
1.27.4 Asymmetric Loss (ASL) (multi-label)	17
1.28 Eigenvectors and Eigenvalues	17
1.29 ALS (Alternating Least Squares)	18
2 Neural Network Training	19
2.1 Batch Norm	19
2.2 Layer Norm	19
2.3 Exploding Gradient	19
2.4 Vanishing Gradient	20
2.5 Gradient Descent	20
2.6 Backpropagation	20
2.7 Adagrad	21
2.8 SGD vs Adam	21
2.9 SGD with Momentum	21
2.10 Adam Optimizer	22
2.11 Weight Decay	22
2.11.1 Weight Decay (Reasoning)	22
2.12 Weight Decay vs L2 Regularization	23
2.13 AdamW	23
2.14 Universal Approximation Theorem	23
2.14.1 wide shallow NN vs. deep narrow NN (M)	23
2.15 Saddle Points	23
3 Transformer and Its Derivatives	24
3.1 Self-Attention	24
3.2 Multi-Head Attention	24
3.3 Decoder	24
3.3.1 Cross-Attention Notes	24
3.3.2 Self-Attention Notes	25
3.4 Inductive Bias	25

3.5	More Attentions	25
3.5.1	Grouped-Query Attention (GQA)	25
3.5.2	Multi-Head Latent Attention (MLA)	25
3.5.3	Linear Attention (Performer)	26
3.5.4	Low-Rank Attention (Linermer)	26
3.6	Attention Collapse	26
3.6.1	Within a single head	26
3.6.2	Across different heads	27
3.7	Activation Functions	27
3.7.1	GELU (Gaussian Error Linear Unit)	27
3.7.2	Swish (a.k.a. SiLU)	27
3.7.3	SwiGLU (Swish-Gated Linear Unit)	27
3.7.4	GLU (Gated Linear Unit)	28
3.8	Positional Embeddings	28
3.8.1	RoPE (Rotary Positional Embedding)	28
3.8.2	NoPE (No Positional Embedding)	28
3.8.3	ALiBi (Attention with Linear Biases)	28
3.8.4	Relative Positional Encoding (RPE)	29
3.8.5	Relative RE (Content + Positional Terms)	29
3.9	More Normalizations	29
3.9.1	RMSNorm (Root mean squared)	29
3.9.2	QKNorm	29
4	Extended Basics	31
4.1	Norms	31
4.2	Linear Independence	31
4.3	Convexity / Concavity	31
4.4	First and Second Order Derivatives	32
4.5	Moments	32
4.6	Probability Distributions	33
4.6.1	Normal Distribution	33
4.6.2	Binomial Distribution	33
4.6.3	Multinomial Distribution	33
4.6.4	Poisson Distribution	33
4.6.5	Beta Distribution	33
4.6.6	Exponential Distribution	34
4.7	Multivariate vs Multimodal Distributions	34
4.8	PDF, PMF, CDF	34
4.9	Independence of Random Variables	34
4.10	Frequentist vs Bayesian	34
4.11	Recommender Systems	34

4.11.1	Collaborative Filtering	34
4.11.2	Assessment	35
4.12	Active Learning	35
4.13	Weakly Supervised Learning	35
4.14	Empirical Risk Minimization (ERM)	36
4.15	Parametric vs Non-parametric Models	36
4.16	Logistic Regression Loss Function	36
4.17	K-means vs GMM	37
4.18	Kernel Methods	37
4.19	Reinforcement Learning (RL)	37
4.19.1	Core Components	37
4.19.2	Exploration vs Exploitation	38
4.19.3	RL Algorithms	38
4.20	TF-IDF	38
5	NLP	39
5.1	Subword Modeling	39
5.2	Assessment Metrics	39
5.2.1	BLEU	39
5.2.2	ROUGE	39
5.2.3	METEOR	40

Part I

ML, Stats, ...

Chapter 1

Basics

1.1 Overfitting vs Underfitting

- Overfitting → high variance → sensitive to noise → poor generalization
- Underfitting → high bias → simple assumptions → poor model

1.1.1 Variance–Bias Tradeoff

- Find “sweet point”
- Solution: model selection, hyperparameter tuning
- In clinic, e.g., @ 90% sensitivity / recall

1.1.2 Solution to overfitting

- Data-level: more training data/data augmentation; data normalisation
- Model-level: simpler models; ensemble methods
- Training-level: drop-out; regularization; weight-decay; batch/layer norm; early stopping; cross-validation

1.2 Precision and Recall

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall (Sensitivity)} = \frac{TP}{TP + TN}$$

$$F_1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

$$\text{Specificity} = \frac{TN}{N}$$

$$\text{FPR} = \frac{FP}{FP + TN}$$

1.3 Ensemble Methods

1.3.1 Bagging:

- Models trained independently (Parallelizable)
- Training data sampled with replacement (bootstrapping)
- Variance reduction
- Example: Random Forest

1.3.2 Boosting:

- Models trained sequentially
- Focus on previous errors / misclassified samples
- Bias reduction
- Example: XGBoost

1.3.3 Stacking Ensemble

- Combine multiple models by training a “meta-learner” to make final decision
- Stage 1: Base models (diverse algorithms: SVM, logistic regression, etc.)
- Stage 2: Meta-model combines predictions of stage 1 models to make final decision (e.g., linear regression, logistic regression)

1.3.4 Cascading Ensemble

- Stage 1: Simple base model. If confidence is high, make decision. If low, move to Stage 2.
- Stage 2: More complex model. Repeat until confidence is high enough to make decision.

1.4 Decision Tree

1. In classification: each leaf is a class.
2. In regression: each leaf is the mean value of all samples belonging to that leaf. i.e., n leaves $\Rightarrow n$ distinct values. We need Random Forest or XGBoost to generate a wide value range (more continuous-like).
3. Not all features are necessarily used, but all might be used.
4. Splitting:
 - Gini impurity / Cross-Entropy for classification.
 - Variance / MAE / MSE for regression.

1.5 Pruning of Decision Trees

Pre-pruning:

- Early stopping: Max depth; Min cluster size; Min decrease of impurity

Post-pruning:

- Grow full tree, then remove nodes that bring little gain by calculating cost for subtrees.

1.6 Entropy, Gini Impurity, and Cross Entropy

Entropy:

$$H = - \sum_i p_i \log(p_i)$$

Measures uncertainty/randomness of a probability distribution.

Gini Impurity:

$$G = 1 - \sum_i p_i^2$$

Cross Entropy:

$$H(p, q) = - \sum_i p_i \log(q_i)$$

Measures the difference between two probability distributions. Also: $H(p, q) = H(p) + KL(p\|q)$

1.7 Random Forest

1. Bagging: sampling with replacement.
2. At each split, only a random subset of features is used to introduce randomness.
This reduces correlation between trees and improves generalisation.

1.8 XGBoost vs LightGBM

1.8.1 XGBoost

Level-wise growth → balanced tree.

- Pros: More stable (especially with small datasets), less prone to overfitting.
- Cons: Slower on large datasets since it considers all splits at each level. May require more splits to reach same loss.

1.8.2 LightGBM

Leaf-wise growth → grows the leaf with the largest potential loss reduction (uneven tree).

- Pros: Lower loss with fewer splits, faster convergence.
- Cons: More prone to overfitting. Requires careful tuning (max depth, min leaf, etc.).

1.8.3 Comparison of Models

	Decision Tree	Random Forest	XGBoost
Structure	Single tree	Bagging / Bootstrapping (Ensemble)	Gradient boosting (Ensemble)
Overfitting	Prone to overfit	Not prone	Avoided by regularization & early stopping
Accuracy	Low accuracy	Better	Best
Interpretability	High	Worse	Worst

1.9 Logistic Regression

$$y = \frac{1}{1 + e^{-(wx+b)}}$$

Sigmoid function. Linear with log-odds.

1.10 Logistic Regression Assumptions

1. Linearity between X and log-odds.
2. Independent observations.
3. No perfect multi-collinearity.

1.11 Linear Regression Assumptions

1. Linearity between X and Y .
2. Independence of errors.
3. Errors have mean 0.
4. Errors have constant variance.
5. No perfect multi-collinearity in features. e.g. $X_3 = \alpha X_1 + \beta X_2$ is perfect multi-collinearity

2,3,4 indicate that **errors are i.i.d..**

1.12 K-means Assumptions

1. Spherical clusters (round, equally sized).
2. Similar cluster sizes.
3. All features contribute equally (scaling needed).
4. Clusters have uniform density.
5. Convex clusters, separable in Euclidean space.

1.13 Feature Selection

- Filter methods: correlation, chi-square, mutual information.
- Wrapper methods: forward selection, backward elimination.
- Embedded methods: L1/L2 regularization, tree-based methods.

1.14 Generative vs Discriminative Models

Generative:

- Learn the underlying probability distribution of the data X and Y .
- Model $P(X, Y)$ or $P(X|Y)$
- Not only classify, but also generate new data
- Examples: GAN, VAE, Naïve Bayes

Discriminative:

- Learn the decision boundary between different classes
- Model conditional probability $P(Y|X)$
- Examples: SVM, KNN

1.15 KL Divergence

- Measures how a probability distribution differs from a reference probability distribution
- Example: In VAE, acts as a regularization term in the loss. Encourages the learned latent distribution to match a prior (normally standard normal).

1.16 L1 vs L2 Regularization

- L1 (Lasso): drives weights to 0. Geometrically in the loss landscape, it forms diamond contour loss with corners.
- L2 (Ridge): small weights but not 0. Geometrically, the loss landscape is elliptical/circular.
- Elastic Net = L1 + L2

1.17 L2-norm vs Squared L2-norm

- Simpler derivative.
- Smooth and differentiable everywhere, including at 0.
- Numerically stable (no division by 0).
- Squared L2-norm is often used in NN

1.18 Likelihood & Posterior

Likelihood: $P(\text{data}|\text{params})$

Posterior: $P(\text{params}|\text{data})$

1.19 Concept Drift (sys design)

- Relationship between X and Y changes over time

1.20 Semi-Supervised Learning

- Example process:
 1. Train supervised model on labeled data
 2. Predict labels for unlabeled data (select high-confidence labels)
 3. Retrain model including pseudo-labels
- Alternatively: joint/iterative training

$$L = L_{\text{sup}} + \lambda L_{\text{unsup}}$$

where L_{sup} = cross-entropy, L_{unsup} = consistency loss / pseudo-label loss.

1.21 Bayes' Rule

$$P(Y | X) = \frac{P(X | Y) \cdot P(Y)}{P(X)}$$

- Prior: $P(Y)$
- Likelihood: $P(X | Y)$
- Evidence: $P(X)$

1.22 Naïve Bayes

- Assumes all features are conditionally independent given the class label.

$$P(x \mid c = c_i) = \prod_{j=1}^N P(x_j \mid c = c_i)$$

- Joint likelihood = product of individual feature likelihoods
- Uses Bayes' rule to compute the posterior of each class given the features
- $P(X, Y) = P(Y) \times P(X|Y)$
- $P(Y \mid X) \propto P(X, Y)$

1.22.1 Naive Bayes Assumptions

All features are conditionally independent given the class.

1.23 AUROC (Receiver Operating Characteristic)

It measures how well a **binary classifier** can distinguish between positive and negative classes. AUROC quantifies a model's ability to rank positives higher than negatives.

$$\begin{aligned} TPR &= \frac{TP}{TP + FN} \\ FPR &= \frac{FP}{FP + TN} \end{aligned}$$

Plot TPR vs FPR.

1.24 R^2 vs Adjusted R^2

- R^2 : monotonic, measures amount of variance explained.
- Adjusted R^2 : penalizes for unnecessary features, accounts for number of features.

1.25 Correlation and Covariance

Covariance:

$$\text{cov}(X, Y) = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})$$

Measures how X and Y vary together.

Correlation:

$$\text{corr}(X, Y) = \frac{\text{cov}(X, Y)}{\sigma_X \sigma_Y}$$

Measures the strength of linear relationship between X and Y .

1.26 Maximum Likelihood Estimation (MLE)

- Estimate parameters of a statistical model.
- Assumes data are from a certain distribution with unknown parameters.
- MLE finds the parameter values that maximize the likelihood of the observed data.

$$L(\theta) = \prod_{i=1}^n P(x_i | \theta)$$

$$\hat{\theta} = \arg \max_{\theta} L(\theta) = \arg \max_{\theta} \sum_{i=1}^n \log P(x_i | \theta)$$

Often easier to compute log-likelihood (sum instead of product), more numerically stable.

1.27 Class Imbalance Loss

1.27.1 Re-weighted Cross Entropy:

$$L = -w_y \log \frac{e^{z_y}}{\sum_j e^{z_j}}$$

where $w_y = \frac{1}{n_c}$, count in class c .

- Pros: handles imbalance
- Cons: gradient instability for very rare classes, may underfit major classes

1.27.2 Balanced Softmax Cross Entropy (multi-class)

$$L = -\log \frac{n_y e^{z_y}}{\sum_j n_j e^{z_j}}$$

- Rebalances within the softmax
- Effective on long-tail distributions, better than re-weighted CE

1.27.3 Focal Loss (binary, multi-class, multi-label)

$$L = -\alpha(1 - p_t)^\gamma \log(p_t)$$

- γ down-weights easy samples; α are class weights
- $\gamma = 0 \Rightarrow$ standard CE
- Larger γ : more imbalance handling (e.g., RetinaNet $\gamma = 2, \alpha = 0.25$ positive, 0.75 negative)

1.27.4 Asymmetric Loss (ASL) (multi-label)

$$L = -y(1 - p)^{\gamma^+} \log(p) - (1 - y)p^{\gamma^-} \log(1 - p)$$

- p : probability from sigmoid, $y \in \{0, 1\}$ ground truth
- Typically $\gamma^+ < \gamma^-$, with $\gamma^+ \in [0, 1]$, $\gamma^- \in [2, 4]$
- Suppresses easy negatives, focuses on hard positives

Comparison:

Focal Loss	ASL
Focus on Hard positives & negatives	Moderate focus on positives, strong suppress negatives
Negatives may dominate loss, positives don't get enough gradient	Works best in multi-label detection with many negatives

1.28 Eigenvectors and Eigenvalues

Given a **square matrix** A , an eigenvector v and its corresponding eigenvalue λ satisfy:

$$Av = \lambda v$$

Eigenvectors show the principal directions of a transformation, and eigenvalues show how strongly those directions are scaled.

1.29 ALS (Alternating Least Squares)

$$M \approx UV^T$$

- Matrix factorization: lower-rank U, V
- Fix U , update V , then fix V , update U :

$$U_i = (V^T V)^{-1} V^T m_i$$

- If $V^T V$ not invertible due to sparsity, use $V^T V + \lambda I$ to ensure positive definiteness

Chapter 2

Neural Network Training

2.1 Batch Norm

- Normalizes activations using statistics computed across the batch dimension.
- Depends on batch size.
- Stabilizes gradient flow
- Allows larger learning rates
- Acts as a form of regularization

2.2 Layer Norm

- Normalizes across hidden dimensions within a single sample, not across the batch
- Independent of batch size
- Work with very small batch, ideal for LLM/Transformers
- Doesn't provide regularization like BatchNorm.

2.3 Exploding Gradient

- Very large gradients cause loss to become NaN or ∞ .

Solutions:

1. Gradient clipping: limit the norm of gradients.
2. Weight initialization:

- Xavier for `tanh/sigmoid`.
 - He for `ReLU`.
3. Normalization: Batch or Layer normalization.
 4. Skip connections: help mitigate, but do not fully solve the problem.

2.4 Vanishing Gradient

- Gradients $\rightarrow 0$, causing slow or stalled training.

Solutions:

1. Use `ReLU` to replace `sigmoid/tanh`.
2. Skip connections (very effective).
3. Normalization.
4. Proper initialization.

2.5 Gradient Descent

- Iteratively update model parameters to minimize loss

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} J(\theta)$$

where $J(\theta)$ = loss function, ∇_{θ} = gradient.

2.6 Backpropagation

- Compute gradient of loss function w.r.t each weight parameter using the chain rule
- Propagate errors backward
- Update parameters via gradient descent (negative gradient direction)

2.7 Adagrad

$$\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,i}} + \epsilon} g_{t,i}$$

$$G_{t,i} = \sum_{i=1}^t g_{t,i}^2$$

- Each parameter θ has its own adaptive learning rate, scaled by the sum of squared gradients.
- If historical g is large, learning rate decreases; if small, learning rate increases.
- Helps with sparse gradients (e.g., in recommender systems or one-hot features where many parameters are infrequently updated).
- Keeps learning rate high and learn efficiently.

2.8 SGD vs Adam

SGD:

- Fixed learning rate
- Can get stuck in local minima
- 1st order of gradients

Adam:

- Adaptive learning rate based on mean and moments
- 2nd order of gradients

2.9 SGD with Momentum

$$\theta_{t+1} \leftarrow \theta_t - \eta \nabla_{\theta} L(\theta_t) + \alpha \Delta \theta_{t-1}$$

Momentum (the last term) introduces an exponential moving average of past gradients.

- Converges more slowly compared to Adam.
- Generalizes better.
- Requires tuning of learning rate.

2.10 Adam Optimizer

$$\begin{aligned}
 m_\theta^t &= \beta_1 m_\theta^{t-1} + (1 - \beta_1) \nabla_\theta L^{t-1} && \text{(1st order moment)} \\
 v_\theta^t &= \beta_2 v_\theta^{t-1} + (1 - \beta_2) (\nabla_\theta L^{t-1})^2 && \text{(2nd order moment)} \\
 \hat{m}_\theta^t &= \frac{m_\theta^t}{1 - \beta_1^t}, \quad \hat{v}_\theta^t = \frac{v_\theta^t}{1 - \beta_2^t} \\
 \theta^t &\leftarrow \theta^{t-1} - \eta \frac{\hat{m}_\theta^t}{\sqrt{\hat{v}_\theta^t} + \epsilon}
 \end{aligned}$$

Typical values: $\beta_1 \approx 0.9$, $\beta_2 \approx 0.999$.

- Combines momentum and adaptive learning rate per parameter.
- Converges faster.
- Generalizes worse (sharp local minima).

2.11 Weight Decay

$$\theta \leftarrow (1 - \lambda)\theta - \eta \nabla_\theta L(\theta)$$

- λ is a small decay factor, e.g., $1e^{-4}$.
- Equivalent to L2 regularization.
- Improves generalization.
- Reduces exploding gradients.
- Biases towards simpler solutions.

2.11.1 Weight Decay (Reasoning)

- We prefer small weights in NN training.
- Large weights amplify sensitivity — small perturbations lead to large changes in output.
- Sensitive to noise and prone to overfitting.
- Creates sharp decision boundaries — fits training data well but generalizes poorly.
- If decay factor is too large $(1 - \lambda) \rightarrow 0$, NN loses learning capacity \Rightarrow underfitting.

2.12 Weight Decay vs L2 Regularization

- In SGD, they are equivalent.
- In Adam, weight decay \neq L2 regularization, since each gradient is rescaled by historical squared gradients.
- Adam has **NO true weight decay**.

2.13 AdamW

Decouples weight decay from gradient update — same principle as SGD.

2.14 Universal Approximation Theorem

Any continuous function can be approximated by a single wide layer NN, but number of neurons may be impractically large.

2.14.1 wide shallow NN vs. deep narrow NN (M)

Certain classes of functions can be represented by deep NN exponentially more compactly than shallow wide NN. Given the same number of neurons, Deep shallow NNs have higher expressiveness and converge faster.

2.15 Saddle Points

At a saddle point, $\nabla L(\theta^*) = 0$, but it is not a local max or min. Hessian has both positive and negative eigenvalues.

- High-dimensional NNs: exponentially more saddle points than local minima.
- Local minima are not problematic in large NNs.
- Saddle points dominate loss landscape \rightarrow main obstacles for gradient descent.
- Optimization gets easier in high dimensions (most local minima are low-loss).
- SGD and momentum help escape saddle points.

Chapter 3

Transformer and Its Derivatives

3.1 Self-Attention

$$\text{Att}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

3.2 Multi-Head Attention

- Multiple self-attention layers in parallel with different learned Q, K, V .
- Allows model to capture information from different representation subspaces.
- Enables focus on different parts of a sequence simultaneously.

3.3 Decoder

1. Masked multi-head self-attention

Uses a causal mask to enforce autoregressive generation. When predicting the next token t , only tokens up to t ($\leq t$) are visible.

2. Encoder–decoder cross-attention

Keys and values come from encoder outputs; queries come from the decoder’s previous layer.

Allows the decoder to condition on the encoded input sequence.

3. Position-wise feed-forward network (FFN)

3.3.1 Cross-Attention Notes

- Happens after masked self-attention and before next-token generation.

- The self-attention hidden states serve as queries.
- During training, cross-attention is applied at every decoded token.
- During inference, the decoder computes cross-attention at time t using only the previously generated tokens, and that cross-attention output is used to generate token t .

3.3.2 Self-Attention Notes

- Masked self-attention during training is computed in parallel.
- During inference it is computed sequentially.

3.4 Inductive Bias

- **CNNs**: Locality and translation invariance.
- **RNNs**: Sequential and temporal inductive bias.
- **Transformers**: Global and context-adaptive; Do not assume temporal or local structure; Serve as general-purpose sequence models.

3.5 More Attentions

3.5.1 Grouped-Query Attention (GQA)

- Share key-value heads for multiple query heads.
- Lower memory, improved efficiency.
- Used in Llama, Qwen, GPT, ...

3.5.2 Multi-Head Latent Attention (MLA)

- Compress key & value heads into lower-dim space (before KV cache)
- Extra matrix multiplication, reduces memory usage.
- Used in DeepSeek.

3.5.3 Linear Attention (Performer)

- Approximates softmax using a kernel feature map ϕ .

$$\text{Att}(Q, K, V) \approx \phi(Q) (\phi(K)^\top V)$$

- Pros: $O(n)$ complexity.
- Cons: approximation error may accumulate.

3.5.4 Low-Rank Attention (LInformer)

- QK^\top is often approximately low-rank.
- Introduce low-rank projections:

$$K' = KE, \quad V' = VE$$

where $\text{rank}(K') \ll \text{rank}(K)$.

$$\text{Att}(Q, K, V) = \text{softmax}(QK'^\top) V'$$

However, these Linear Attention and Low-Rank Attention mechanisms are rarely used in LLMs due to their approximation errors.

3.6 Attention Collapse

- Attention weights across tokens become uniform or overly concentrated.

3.6.1 Within a single head

Examples:

$$[0.2, 0.2, 0.2, 0.2, \dots] \quad \text{or} \quad [1, 0, 0, 0, 0, \dots].$$

Solution:

- Proper initialization
- Scaled dot-product
- Pre-layer normalization

3.6.2 Across different heads

- Heads learn similar weights or become highly correlated.

Solution:

- Head dropout
- Head diversity regularization

3.7 Activation Functions

3.7.1 GELU (Gaussian Error Linear Unit)

$$\text{GELU}(x) = x \cdot \Phi(x)$$

where $\Phi(x)$ is the CDF of the standard normal distribution.

- Smooth, improves convergence compared to ReLU in deep transformers.
- Used in Transformer, BERT, GPT-2.
- Original activation function in Transformer.
- Computationally more expensive than SiLU, less used now.

3.7.2 Swish (a.k.a. SiLU)

$$f(x) = x \cdot \sigma(x)$$

- Similar to GELU, smoother decay for small negative values.
- Used in EfficientNet, NLP models.

3.7.3 SwiGLU (Swish-Gated Linear Unit)

$$f(x) = (XW_1) \odot \text{Swish}(XW_2)$$

- Combines GLU with Swish (element-wise gating).
- Lets the network learn when to pass signals, improves expressiveness.
- Helps gradient flow and non-linear expressivity.
- Used in LLaMA, PaLM, GPT-4.

3.7.4 GLU (Gated Linear Unit)

$$f(x) = (XW_1) \odot \sigma(XW_2)$$

- Splits into content and gate.
- σ is sigmoid activation.

3.8 Positional Embeddings

- Transformers have no recurrence or convolution, so they need a way to encode order of tokens.
- The (original) positional embeddings (a.k.a. **absolute positional embeddings**) are added to the input embeddings.

3.8.1 RoPE (Rotary Positional Embedding)

- Rotates Q and K vectors relative to their token position.
- Encodes position implicitly by rotating angles.
- No need for extra embedding parameters.
- Lightweight, works better in long sequences.
- Used in LLaMA, Qwen, ChatGLM.

3.8.2 NoPE (No Positional Embedding)

- Typically combined with RoPE for better length generalization.
- Aware of position because of causal attention mask.

3.8.3 ALiBi (Attention with Linear Biases)

$$\text{Attn}(i, j) = \frac{Q_i K_j^T}{\sqrt{d}} + m_h(i - j)$$

- Adds linear bias term depending on relative distance.
- Encourages attention to focus on nearby tokens.
- Simple and cheap computation.
- Used in LLaMA-2, MPT.

3.8.4 Relative Positional Encoding (RPE)

$$\text{Attn}(i, j) = \frac{Q_i K_j^T}{\sqrt{d}} + b_{ij}$$

- b_{ij} = learned embedding/bias depending on relative distance.
- High expressiveness but more expensive to compute.
- Used in Transformer-XL, T5.

3.8.5 Relative RE (Content + Positional Terms)

$$\text{Attn}(Q, K) = QK^T + QR^T + uK^T + vR^T$$

- R encodes relative position.
- Handles long context but memory inefficient.
- Does not generalize well beyond training sequence lengths.

3.9 More Normalizations

3.9.1 RMSNorm (Root mean squared)

$$\text{RMS}(x) = \sqrt{\frac{1}{d} \sum_{i=1}^d x_i^2 + \epsilon}$$

$$\text{RMSNorm}(x) = \frac{x}{\text{RMS}(x)} \odot w$$

- w is a scalar
- Doesn't centre data, doesn't compute variance.
- Less computations, more stable.

3.9.2 QKNorm

- Avoids QK^T gets too large (softmax saturates)
- Attention:

$$\text{softmax} \left(\frac{\tilde{Q}\tilde{K}^T}{\tau} \right) V$$

Where:

$$\tilde{Q} = \text{norm}(Q), \quad \tilde{K} = \text{norm}(K)$$

- $\text{norm}(\cdot)$ could be L2-norm, RMSNorm.
- τ is temperature, a scaling factor.

Chapter 4

Extended Basics

4.1 Norms

Properties of a norm:

1. Non-negativity: $\|u\| \geq 0$, and $\|u\| = 0 \Rightarrow u = 0$
2. Homogeneity: $\|\alpha v\| = |\alpha| \|v\|$
3. Triangle inequality: $\|u + v\| \leq \|u\| + \|v\|$

L_p norm:

$$\|x\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}$$

Special case: L_0 counts non-zero elements.

Note: Norm \Rightarrow Metric, but Metric $\not\Rightarrow$ Norm.

4.2 Linear Independence

For $c_1a + c_2b = 0$, the only solution is $c_1 = c_2 = 0$.

4.3 Convexity / Concavity

- Convex: $f''(x) \geq 0, \forall x$
- Concave: $f''(x) \leq 0, \forall x$

4.4 First and Second Order Derivatives

Taylor expansion:

$$L(w + \Delta w) \approx L(w) + \nabla L(w)^T \Delta w + \frac{1}{2} \Delta w^T H \Delta w$$

where $H = \nabla^2 L(w)$ is the Hessian matrix.

Update rule:

$$\Delta w = -H^{-1} \nabla L(w)$$

First order ($\nabla L(w)$): set slope

Second order (H): set step size

- Could help with adaptive gradient descent, however due to its memory and computation complexity, rarely used in NN!

4.5 Moments

n -th moment:

$$\mu_n = E[(X - \mu)^n]$$

- 0th: total probability = 1
- 1st: mean
- 2nd: variance
- 3rd: skewness
- 4th: kurtosis

Note: Adam and RMSProp use 2nd order moments, NOT Hessian!

1st order: mean of gradients.

2nd order: variance of gradients, helps smooth noisy gradients, handle parameter scales.

- When 2nd order moment is large \Rightarrow more uncertainty \Rightarrow smaller step size.

4.6 Probability Distributions

4.6.1 Normal Distribution

$$X \sim \mathcal{N}(\mu, \sigma^2)$$

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

$$E[X] = \mu, \quad Var(X) = \sigma^2$$

4.6.2 Binomial Distribution

$$X \sim Bin(n, p)$$

$$P(X = k) = \binom{n}{k} p^k (1-p)^{n-k}, \quad E[X] = np, \quad Var(X) = np(1-p)$$

4.6.3 Multinomial Distribution

$$X \sim Multi(n, \pi)$$

where $\pi = (p_1, p_2, \dots, p_k)$, $\sum p_i = 1$.

$$P(X_1 = x_1, \dots, X_k = x_k) = \frac{n!}{x_1! \cdots x_k!} \prod p_i^{x_i}, \quad E[X_i] = np_i, \quad Var(X_i) = np_i(1-p_i)$$

4.6.4 Poisson Distribution

$$X \sim Poisson(\lambda)$$

$$P(X = k) = \frac{\lambda^k}{k!} e^{-\lambda}, \quad E[X] = \lambda, \quad Var(X) = \lambda$$

Poisson \approx Binomial when n large and p small.

4.6.5 Beta Distribution

$$X \sim Beta(\alpha, \beta), \quad \alpha, \beta > 0$$

$$f(x; \alpha, \beta) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1-x)^{\beta-1}, \quad E[X] = \frac{\alpha}{\alpha + \beta}, \quad Var(X) = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}$$

4.6.6 Exponential Distribution

$$f(x; \lambda) = \lambda e^{-\lambda x}, \quad E[X] = \frac{1}{\lambda}, \quad Var(X) = \frac{1}{\lambda^2}$$

4.7 Multivariate vs Multimodal Distributions

- Multivariate: multiple variables (e.g., distribution of (X, Y, Z)).
- Multimodal: multiple peaks/modes in PDF (e.g., mixture of two Gaussians).

4.8 PDF, PMF, CDF

- PMF: discrete variable probability at x
- PDF: continuous variable density at x
- CDF: total probability accumulated up to x

4.9 Independence of Random Variables

$$f_{X,Y}(x, y) = f_X(x)f_Y(y)$$

4.10 Frequentist vs Bayesian

	Frequentist	Bayesian
Parameter	Fixed	Random
Probability	Only for data	Both parameter & data
Inference	Point estimate, CI, p-values	Posterior, credible intervals
Prior	None	Yes

4.11 Recommender Systems

4.11.1 Collaborative Filtering

- Recommends items based on user-item interaction history

4.11.2 Assessment

1. **Explicit ratings:** RMSE, MAE, R^2
2. **Implicit feedback:** Precision/Recall@k, NDCG@k, MAP (Mean Average Precision)
3. **Rankings:** NDCG@k, MAP
4. **Business-oriented metrics:** CTR (Click-through rate), Conversion rate

4.12 Active Learning

Model actively selects which data points to label to maximize learning efficiency.

Steps:

1. Train initial model on small labeled set.
2. Model identifies uncertain/informative samples.
3. Query an oracle (human annotator) for labels.
4. Retrain model with new labels → repeat.

Strategies:

- Least confidence: pick samples where top predicted class has lowest probability.
- Margin sampling: pick samples where the difference between top two class probabilities is smallest.
- Entropy-based: pick samples with highest entropy in prediction.

4.13 Weakly Supervised Learning

Labels are incomplete, noisy, bad-quality, or inexact.

- Incomplete label → semi-supervised learning.
- Noisy labels: use robust loss functions.
 - e.g. MAE instead of CE for classification.
 - Label smoothing: soften one-hot labels

$$y'_i = (1 - \epsilon)y_i + \epsilon/k$$

- Generalized CE.
- Sample weighting.
- Soft label loss:

$$y_{target} = \beta y_{noisy} + (1 - \beta) y_{model}$$

4.14 Empirical Risk Minimization (ERM)

Find a model to minimize expected loss on data:

$$R(f) = \mathbb{E}_{(x,y) \sim P_{data}} [L(f(x), y)]$$

$L(\cdot)$ is the loss function, e.g. CE. P_{data} is the true data distribution.

Unknown true data distribution → approximate with empirical average (observed data):

$$\hat{R}(f) = \frac{1}{N} \sum_{i=1}^N L(f(x_i), y_i)$$

⇒ supervised learning task solved via closed form (if possible) or use NN by gradient descent.

4.15 Parametric vs Non-parametric Models

	Parametric	Non-Parametric
Definition	Fixed form for function/distribution	No assumptions on form
Parameters	Fixed #	Grow with data
Flexibility	Less	More (can approx. a wide range of functions)
Computation	Faster	Slower
Examples	Logistic Regression	kNN, k-means, Decision Trees, Gaussian Process

4.16 Logistic Regression Loss Function

Use log-loss (cross entropy) for logistic regression, not MSE.

- MSE causes non-convexity when applied to logistic regression.
- Gradient saturation when $\hat{y} \rightarrow 0/1 \Rightarrow$ slow learning.
- CE is convex ⇒ global minimum guaranteed.

- CE provides strong gradient when predictions are wrong.
- Theoretically optimal under Bernoulli distribution assumption.
- MSE treats probabilities linearly, but predicted probabilities are log/sigmoid transformed → mismatch.

4.17 K-means vs GMM

- K-means: spherical clusters, hard assignment, non-probabilistic, sensitive to outliers, Lloyd's algorithm.
- GMM: elliptical clusters, soft probabilistic assignments, more robust but computationally heavier, EM algorithm.

4.18 Kernel Methods

1. Linear: $k(x, x') = x^T x'$
2. Polynomial: $k(x, x') = (x^T x' + c)^d$
3. RBF/Gaussian: $k(x, x') = \exp(-\|x - x'\|^2 / 2\sigma^2)$
4. Sigmoid: $k(x, x') = \tanh(\alpha x^T x' + \epsilon)$

4.19 Reinforcement Learning (RL)

4.19.1 Core Components

- **Agent:** The learner.
- **Environment:** Everything the agent interacts with.
- **State (s):** Description of the environment at time t .
- **Action (a):** Choice the agent makes given s_t .
- **Reward (r):** Feedback signal indicating the benefit of an action.
- **Policy (π):** The strategy the agent uses to make actions based on states.
- **Value function (V):** Estimates how good it is for the agent to be in a state or take an action, in terms of expected future rewards.

4.19.2 Exploration vs Exploitation

Balance between trying new actions vs choosing the best known action.

4.19.3 RL Algorithms

1. **Value-based:** Q-learning
2. **Policy-based:** PPO, REINFORCE
3. **Actor-Critic:** Combines (1) and (2), e.g., A2C, SAC

4.20 TF-IDF

- Term Frequency - Inverse Document Frequency
- Measures word importance in a document

Chapter 5

NLP

5.1 Subword Modeling

Reduce vocab size → solve OOV (out of vocabulary) problem. Improve data efficiency.

- **Byte-Pair Encoding (BPE)**: adopted by GPT.
- **WordPiece**: used in BERT.
- **Unigram LM Tokenization / SentencePiece**: used in LLaMA, T5.

5.2 Assessment Metrics

5.2.1 BLEU

- Compares n -gram, computes precision of n -gram, applies brevity penalty (BP).

$$\text{BLEU} = \text{BP} \cdot \exp \left(\sum_{i=1}^N w_i \log p_i \right)$$

p_i is precision.

- Usually computed for unigram to 4-gram.
- **Pros**: Standard, fast, automatic (no human), language-independent.
- **Cons**: Insensitive to semantics, does not capture grammar or fluency.

5.2.2 ROUGE

- Recall-oriented.
- Measures overlap between languages (recall focus, unlike BLEU which is precision).

- **ROUGE-N**: n -gram recall.
- **ROUGE-L**: longest common subsequence.
- **ROUGE-S**: skip-bigram overlap.
- Good to measure summarization.

5.2.3 METEOR

- Considers precision, recall, and alignments such as (Step 1):
 - Exact word match.
 - Stem match (e.g., run \leftrightarrow running).
 - Synonyms.
- Step 2: Computes unigram precision and recall.
- Step 3: Computes F1-score.
- Step 4: Fragmentation penalty (penalize word order differences).